

# Facebook's Secure and Private AI Scholarship Challenge 2019

Notes by Pranjal Chaubey

<https://www.linkedin.com/in/pranjallchaubey/>

<https://github.com/pranjalchaubey>



UDACITY

# Deep Learning with PyTorch

5<sup>th</sup> June 2019

Best method for ~~matrix~~<sup>tensor</sup> multiplication in PyTorch

`torch.mm(tensor1, tensor2)`

`torch.matmul()` — Broadcasts tensors, hence could give results even when the dimensions don't match.  
Not Recommended!

Methods to manipulate the shape of ~~tensor~~ tensor

`weights.reshape(a, b)` — 'Sometimes' copies the whole tensor to a new memory  
— NOT EFFICIENT!

`weights.resize_(a, b)` — If the new shape doesn't match, it will add/subtract new elements in the tensor without giving any error.  
— NOT RECOMMENDED!

`weights.view(a, b)` — Does what is asked, in the way it is expected.  
— HIGHLY RECOMMENDED

`.view(-1)` — Flattens a tensor

`images.view(images.shape[0], -1)` — Flattens the 'image' tensor in this case  
Size of the Batch

Loading the training data through DataLoader  
trainloader = torch.utils.data.DataLoader

(trainset, batch\_size=64, shuffle=True)

Training Dataset

When iterating through 'trainloader',  
each iteration will contain a  
batch of 64 images.

Randomly pick images  
in a batch for training

Converting 'trainloader' in an iterable,

dataiter = iter(trainloader)

images, labels = dataiter.next()

Sigmoid Activation Function

def activation(x):

return 1/(1+torch.exp(-x))

Softmax Function

def softmax(x)

return torch.exp(x)/torch.sum(torch.exp(x),  
dim=1).view(-1, 1)

$$\text{Softmax} = \sigma(x_i) = \frac{e^{x_i}}{\sum^K e^{x_k}}$$

# Building Networks with PyTorch nn Module

from torch import nn

class Network(nn.Module):

def \_\_init\_\_(self):

super().\_\_init\_\_()

Registers the 'network' with  
nn module

# Inputs to hidden layer linear transformations  
self.hidden = nn.Linear(784, 256)

# Output Layer

self.output = nn.Linear(256, 10)

# Define sigmoid activation & softmax output

self.sigmoid = nn.Sigmoid()

self.softmax = nn.Softmax(dim=1)

def forward(self, x)

This section can be removed  
using nn.functional

# Pass the input tensor through each operation

x = self.hidden(x)

x = self.sigmoid(x)

x = self.output(x)

x = self.softmax(x)

return x

```

import torch.nn.functional as F
class Network(nn.Module):
    def __init__(self):
        super().__init__()
        # Inputs to hidden layer linear transformation
        self.hidden = nn.Linear(784, 256)
        # Output Layer
        self.output = nn.Linear(256, 10)

    def forward(self, x):
        # Hidden layer with sigmoid activation
        x = F.sigmoid(self.hidden(x))
        # Output Layer with Softmax
        x = F.softmax(self.output(x), dim=1)
        return x

```

A more succinct version of the previous network

```

# Initialize the model by creating the network object
model = Network()

```

## Using nn.Sequential

# Hyperparameters for our network

input\_size = 784

hidden\_sizes = [128, 64]

output\_size = 10

from collections import OrderedDict

model = nn.Sequential(OrderedDict([

Layer  
names  
have to  
be  
UNIQUE

(  
('fc1', nn.Linear(input\_size, hidden\_sizes[0])),  
('relu1', nn.ReLU()),  
('fc2', nn.Linear(hidden\_sizes[0], hidden\_sizes[1])),  
('relu2', nn.ReLU()),  
('output', nn.Linear(hidden\_sizes[1], output\_size)),  
('softmax', nn.Softmax(dim=1))])])

Ultra fast method to define a network

model.fc1.weight } Accessing various model layers  
model.fc1.bias }

model.fc1.bias.data.fill\_(0) } Setting all biases to 0

model.fc1.weight.data.normal\_(std=0.01)

| Setting weights to a sample from random normal distribution with a Standard Deviation of 0.01

# Training a Neural Network

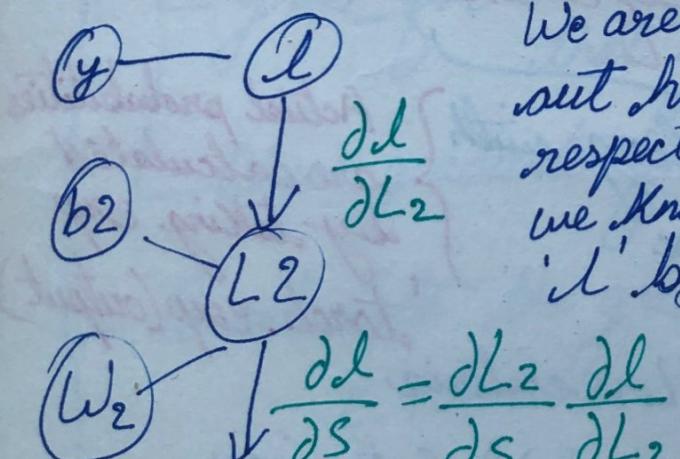
7<sup>th</sup> June 2019

We need a loss function to determine the errors our neural network is making, in case of ~~classification~~ regression and binary ~~classification~~ classification problems, often Mean Squared Error is used.

$$l = \frac{1}{2n} \sum_i^n (y_i - \hat{y}_i)^2$$

Number of Examples      True Labels      Predicted Labels

## Backpropagation



We are effectively trying to figure out how ' $l$ ' is changing with respect to a change in ' $w_1$ '. Once we know that, we can minimize ' $l$ ' by updating weights,

$$w_1 = w_1 - \alpha \frac{\partial l}{\partial w_1}$$

[Learning Rate]

$$\frac{\partial l}{\partial L_1} = \frac{\partial S}{\partial L_1} \frac{\partial l}{\partial S} = \frac{\partial S}{\partial L_1} \frac{\partial L_2}{\partial S} \frac{\partial l}{\partial L_2}$$

$$\frac{\partial l}{\partial w_1} = \frac{\partial L_1}{\partial w_1} \frac{\partial l}{\partial L_1}$$
$$= \frac{\partial L_1}{\partial w_1} \frac{\partial S}{\partial L_1} \frac{\partial L_2}{\partial S} \frac{\partial l}{\partial L_2}$$

Gradient

We typically use the nn. CrossEntropyLoss in PyTorch in classification problems.

Conventionally, the loss is assigned to a variable called criterion.

CrossEntropyLoss requires LOGITS as its inputs instead of the probabilities from the Softmax.

↳ Logits are nothing but the Inputs to the Softmax function.

Cross Entropy Loss is a combination of

nn.NLLLoss

nn.LogSoftmax

Negative Log

Likelihood Loss

Used along with

LogSoftmax

Actual probabilities  
are calculated  
by taking exp.  
'torch.exp(output)'

criterion = nn.NLLLoss() Loss criteria

images, labels = next(iter(trainloader)) Get images

images = images.view(images.shape[0], -1) Flatten images

logits = model(images) Get the logits and make a Forward Pass

loss = criterion(logits, labels) Calculate loss

print(loss)

# AUTOGRAD

Autograd automatically calculates the gradients.

Set `requires_grad=True` for a tensor

Can be turned off for a block of code

$x = \text{torch.zeros}(1, \text{requires\_grad}=\text{True})$

with `torch.no_grad()` — Turns off autograd

$$y = x^2$$

Gradients can be globally turned on and off using `torch.set_grad_enabled(True/False)`

Keeps track of all the operations performed on a tensor

`grad_fn` can be used to check the function that generated this variable

Check gradients for a tensor `tensor.grad`

To calculate the gradient of  $z$  w.r.t.  $x$

`z.backward()`

When a new network is created in PyTorch, all the parameters are initialized with `requires_grad=True`

8<sup>th</sup> June 2019

## OPTIMIZER

Once we get the gradients from Autograd

↳ We need to update the weights

↳ performed by an optimizer

Optimizers in PyTorch live inside the optim package

↳ from torch import optim

optimizer = optim.SGD(model.parameters(), lr=0.01)

Stochastic Gradient  
Descent

Model  
parameters to  
optimize

Learning  
Rate

General Steps to train a network in PyTorch

↳ Zero out the gradients since they keep getting accumulated with each batch

optimizer.zero\_grad()

Make a Forward Pass through the network

Use the network output to calculate the loss

Perform backward pass using loss.backward() to calculate the gradients

Take a STEP with the optimizer to update the weights

One pass through the entire dataset is called an EPOCH

```
from torch import optim
optimizer = optim.SGD(model.parameters(), lr=0.01)
images, labels = next(iter(trainloader))
images.resize_(64, 784)
optimizer.zero_grad()
output = model.forward(images)
loss = criterion(output, labels)
loss.backward()
optimizer.step()
```

-----  
epochs = 5

for e in range(epochs):

running\_loss = 0

for images, labels in trainloader:

images = images.view(images.shape[0], -1)

optimizer.zero\_grad()

output = model.forward(images)

loss = criterion(output, labels)

loss.backward()

optimizer.step()

running\_loss += loss.item()

else:

print(f"Training loss {running\_loss / len(trainloader)}")

TRAINING Loop

10<sup>th</sup> June 2019

## VALIDATION

After training the model on 'training data', we need to 'validate' it on 'testing data' in order to check its accuracy.

Training data is loaded in trainloader.  
Testing data is loaded in testloader.

Finding the model accuracy on a test batch is simple,

# Get the class probabilities

$ps = \text{torch.exp}(\text{model.forward(images)})$

# Use  $ps.\text{topk}(n)$  method that gives the highest 'n' classes predicted for an input

#  $ps.\text{topk}(1)$  gives the ~~highest~~ class with the highest probability

~~top\_p, top\_class = ps.topk(1, dim=1)~~

~~MARK~~ # Get the correctly predicted labels as 1

~~equals = top\_class == labels.view(\*top\_class)~~

To ensure correct dimensions of tensors

~~ByteTensor~~

# Get the accuracy

~~accuracy = torch.mean>equals.type(torch.FloatTensor))~~

Connect to float to calculate mean →

for e in range (epochs):

    for images, labels in trainloader:

        } Training loop

else:

    test\_loss = 0

    accuracy = 0

    with torch.no\_grad():

        for images, labels in testloader:

            log\_ps = model.forward(images)

            test\_loss += criterion(log\_ps, labels)

            ps = torch.exp(log\_ps)

            top\_p, top\_class = ps.topk(1, dim=1)

            equals = top\_class == labels.view(\*top\_class)

            accuracy += torch.mean(equals.type(torch.FloatTensor))

        print(f'Test loss: {test\_loss / len(testloader)}')

        print(f'Accuracy: {accuracy / len(testloader)}')

VALIDATION Loop

## OVERFITTING

As the network tends to learn the training data better and better, it tends to **OVERFIT**.

As a result, its accuracy on the **VALIDATION / TESTING DATA** goes **DOWN!**

To avoid overfitting

Early Stopping - Save different models and choose the one with lowest validation loss.

**Dropout** - Randomly drop neural units to make the network more robust.  
ref. dropout = nn.Dropout( $p=0.2$ )

$x = \text{self.dropout}(\text{F.relu}(\text{self.fc1}(x)))$

During testing/validation we would want to use the entire model without dropout.

model.eval()

Back to training mode

model.train()

## SAVING and LOADING MODELS

The model parameters for PyTorch networks are stored in a model's state\_dict.

↳ `print(model.state_dict().keys())`

Saving the models is easy

↳ `torch.save(model.state_dict(), 'filename.pth')`

Extension for  
PyTorch Models

Loading the state\_dict

↳ `state_dict = torch.load('filename.pth')`

Attaching the state\_dict to a new model

↳ `model.load_state_dict(state_dict)`

The new model needs to have exactly the same number of layers and the number of neurons per layer as the original ~~model~~ model. Otherwise, PyTorch will throw an error.

## LOADING IMAGE DATA

Simplest way to load image data in PyTorch is through datasets. `ImageFolder` from `torchvision`

`dataset = datasets.ImageFolder('path/to/data', transform=transform)`

The system expects the folder to have a particular structure,

`root/dog/d1.png`

`root/dog/d2.png`

`root/cat/c1.png`

`root/cat/c2.png`

⋮

Images are also usually transformed when they are loaded

`transform = transforms.Compose([`

Transforms

Pipeline

`transforms.Resize(255),  
transforms.CenterCrop(224),  
transforms.ToTensor()])`

We generally introduce randomness into the data used for training the network to make it more robust

`RandomRotation(30)`

`RandomResizedCrop(224)`

`RandomHorizontalFlip()`

`Normalize([0.5, 0.5, 0.5], [0.5, 0.5, 0.5])`

Keeps the network weights near 0

Makes backpropagation more stable

## TRANSFER LEARNING

We use pre-trained models to solve hard Computer Vision problems

[ from torchvision import datasets, transforms, model  
imports pre-trained models ]

Loading the models model=models.densenet121  
(pretrained=True)

Need to freeze the model parameters so that we don't backpropagate through them

for param in model.parameters():

param.requires\_grad=False

We only modify the final classifier layer according to our applications

classifier = nn.Sequential( : )  
New layers that we will attach as the first layers in the pretrained network

model.classifier = classifier } Attach the new layers to the pretrained model

Moving tensors back and forth from GPU to CPU

model.cuda() model.cpu()

images.cuda() images.cpu()

print(torch.cuda.is\_available()) Check if cuda is available

model.to('cuda')

model.to('cpu')